

Sound and Partially-Complete Static Analysis of Data-races in GPU Programs

Dennis Liew¹, Tiago Cogumbreiro¹, Julien Lange²

¹University of Massachusetts Boston, USA. ²Royal Holloway, University of London, UK

Introduction

GPUs are parallel computation devices with high susceptibility for bugs such as data-races which may produce non-deterministic behaviour.

We proposed *approximation analysis*, a static analysis technique to detect true data-races in GPU kernels, which tests whether accesses are **reachable (Control Independence)**, and when the reported locations are **precise (Data Independence)**.

Our theory was implemented in the tool FaialAA, as the first sound and partially-complete DRF verifier that can detect true data-races.



When are reports from bug finding tools true?

In our evaluation (3 datasets) Faial:

- Reported **2.1× fewer** potential alarms in a dataset of 227 kernels
- Found **12 undocumented** racy kernels, including 8 that are missed by GPUVerify and Faial from a well-studied dataset.
- Certified **5 documented bugs** (OpenMM and Nvidia’s Megatron-LM) **and their fixes**, while others only succeed in 2 documented bugs+fixes.

Faial is a static analyzer for CUDA kernels that can check for racy and data-race free kernels.
Try our GitHub Action!



Evaluation

Kernel	GPUVerify		Faial		Faial + AA	
	Racy	DRF	Racy	DRF	Racy	DRF
bucketPos	P-R	DRF	P-R	DRF	T-R	DRF
compRange	P-R	DRF	P-R	DRF	T-R	DRF
reduceVal	P-R	✖	n/a	n/a	T-R	DRF
sortBucket	t/o	t/o	n/a	n/a	P-R	✖
gradInput	✖	DRF	n/a	n/a	T-R	DRF
layerNorm	✖	DRF	n/a	n/a	T-R	DRF

Tools	Kernels				Alarms	
	DRF	P-Racy	T-Racy	Unsupported	True	Potential
GPUVerify	193	17	n/a	16	n/a	50
Faial	207	11	n/a	8	n/a	21
FaialAA	210	4	12	0	22	10

Theoretical Results

Theorem 4.5 (True Positives): identifies a specific class of programs where our analysis only report **true alarms**.

Let $\emptyset \vdash CI \llbracket [s] \rrbracket$, $\emptyset \vdash DI \llbracket [s] \rrbracket$, and $datarace(\delta_1, \delta_2)$.

If $\delta_1 \in p\text{-actions}(\llbracket [s] \rrbracket)$ and $\delta_2 \in p\text{-actions}(\llbracket [s] \rrbracket)$, then $\delta_1 \in j\text{-actions}(s)$ and $\delta_2 \in j\text{-actions}(s)$.